

# Reproducing GOAL in 3D Torus

## AI+X - Lab4

Ziren Wang, Chuan Liu

Yao Class 23 & 22

# I. Preliminaries

# 3D Torus

## Definition

A Torus, or  $k$ -ary  $n$ -cube, is an  $n$  dimensional torus network with  $k$  nodes per dimension giving a total of  $N = k^n$  nodes. It has following properties and therefore becomes popular in network routing algorithm research:

- $\mathcal{O}(N)$  cost
- Exploit locality for near-neighbor traffic
- High path diversity
- Edge symmetric

# Routing algorithm

## Definition

A routing algorithm  $R$  maps a source-destination pair to a path through the network from the source to the destination. We can divide routing algorithms by the following criteria:

- **Oblivious** algorithms select the path using only the identity of the source and destination nodes.
- **Adaptive** algorithms may also base routing decisions on the state of the network.
- **Minimal** algorithms only route packets along a shortest path from source to destination.
- **Non-minimal** ones may route packets along longer paths.

## II. GOAL

# Algorithm details

The GOAL algorithm contains two parts:

---

## Algorithm 1 Selecting quadrant

---

**Input:** Source node  $s$ , destination node  $d$

**Output:** direction vector  $r$

$x_s, y_s, z_s \leftarrow$  coordinates of  $s$

$x_d, y_d, z_d \leftarrow$  coordinates of  $d$

$\delta_x = \min(x_d - x_s, x_s - x_d)$

$p_x^+ = 1 - \delta_x/k$

$r_x \sim \text{Bernoulli}(p_x^+)$

$r_y$  and  $r_z$  are updated similarly

---

---

## Algorithm 2 Adaptive routing

---

**Input:**  $s, d, r$

**Output:** path from  $s$  to  $d$

$x_d, y_d, z_d \leftarrow$  coordinates of  $d$

$x, y, z \leftarrow x_s, y_s, z_s$

$path \leftarrow \emptyset$

**while**  $(x, y, z) \neq (x_d, y_d, z_d)$  **do**

    calculate productive dimensions

    choose shorter queue length

$path \leftarrow path + \text{next hop}$

**end while**

**return**  $path$

---

# Flow Control

## Virtual Channels

Our implementation of GOAL employs 4 VCs per physical channel to achieve deadlock freedom in the network. We group the 4 VCs into 3 pairs,  $*_0$ ,  $*_1$ , and non- $*$ .

- non- $*$  channels are fully adaptive and can be used at any time.
- $*$  can be used only when the packets are traversing the *most significant* productive dimension.
- Packets will (will not) move through  $*_1$  ( $*_0$ ) when crossing a wrap-around edge.

### III. Implementation



# Routing

## Adaptive Routing

For Goal routing algorithm, we implement `OutputUnit::count_free_vc`. That's why GOAL routing can find the most `vc_free` outputport and chose it.

```
40 namespace garnet
115 OutputUnit::has_free_vc(int vnet, bool is_ring, bool is_ring_checkpoint, bool is_torus, std::vector<bool> is_torus_dlas_checkpoint, bool is_minimal_torus, Routing
118 int found = get_dla_of_direction();
119 if (is_ring) {
120     if (is_ring_checkpoint) {
121         for (int vc = vc_base; vc < vc_base + m_vc_per_vnet/2; vc++) {
122             if (is_vc_idle(vc, curTick()))
123                 return true;
124         }
125     }
126     else {
127         for (int vc = vc_base + m_vc_per_vnet/2; vc < vc_base + m_vc_per_vnet; vc++) {
128             if (is_vc_idle(vc, curTick()))
129                 return true;
130         }
131     }
132 }
133 else if (is_torus) {
134     if (found < 0) {
135         for (int vc = vc_base; vc < vc_base + m_vc_per_vnet; vc++) {
136             if (is_vc_idle(vc, curTick()))
137                 return true;
138         }
139     }
140     if (routing_algorithm == Goal || routing_algorithm == Min_AD) {
141         for (int vc = vc_base; vc < vc_base + m_vc_per_vnet/2; vc++) {
142             if (is_vc_idle(vc, curTick()))
143                 return true;
144         }
145     }
146     if (is_torus_dlas_checkpoint[found]) {
147         for (int vc = vc_base + m_vc_per_vnet/2; vc < vc_base + m_vc_per_vnet; vc++) {
148             if (is_vc_idle(vc, curTick()))
149                 return true;
150         }
151     }
152     else {
153         for (int vc = vc_base + m_vc_per_vnet/2; vc < vc_base + m_vc_per_vnet; vc++) {
154             if (is_vc_idle(vc, curTick()))
155                 return true;
156         }
157     }
158 }
159 }
160 }
```

# Flow Control

## Virtual Channels

We maintain an `std::vector<bool>` in `RouteInfo`, each element assigned with a dimension and set to be true when the flit goes through a wrap-around edge.

Before send the flit, GOAL will call `route_compute` again and calculate new output in time. This operation avoids deadlock.

```
41 namespace goal {
42     bool route_compute(CircuitInfo& info, RouteInfo& route,
43                       std::vector<bool> wrap_around) {
44         // Create a random number from hardware
45         std::mt19937 gen(info); // Seed the generator
46         std::uniform_int_distribution<int> di(0, 100);
47
48         route.directions.resize(wrap_around.size()); // Resize to hold wrap_around elements
49         for (int i = 0; i < wrap_around.size(); ++i) {
50             if (gen_probability() < 1) (route.directions[i] += 1) % 4; continue;
51             double rand_num = di(gen);
52             route.directions[i] = (rand_num < gen_probability()) ? 1 : -1;
53
54             // If (gen_probability() < 0.5) route.directions[i] = 1;
55             // else route.directions[i] = -1; // for testing
56
57             if (gen_probability() < 0.5 && route.directions[i] == 1) (route.is_minimal_turn = true);
58             else if (gen_probability() < 0.5 && route.directions[i] == -1) (route.is_maximal_turn = true);
59         }
60
61         std::vector<bool> check_output_info;
62         for (int i = 0; i < wrap_around.size(); ++i) {
63             if (route.directions[i] == 0) continue;
64             route.directions[i] = 0;
65             continue;
66         }
67         for (int i = 0; i < wrap_around.size(); ++i) {
68             if (route.directions[i] == 1) {
69                 temp_output = "left" + std::to_string(i) + "_port";
70             } else if (route.directions[i] == -1) {
71                 temp_output = "right" + std::to_string(i) + "_port";
72             }
73             check_output_info.push_back((info.outputs[i] == 0) && temp_output);
74         }
75
76         // Check these outputs, count their vcs and pick the one with minimal occupation in vc.
77         int min_vc_free = -1;
78         int output = check_output_info[0];
79         bool is_output_free = routing_algorithm;
80         for (int i = 0; i < check_output_info.size(); ++i) {
81             int vc_free = output_vc_free - count_free_vc(route_vcs, false, false, true, route.is_minimal_turn, routing_algorithm);
82             if (vc_free > min_vc_free) {
83                 min_vc_free = vc_free;
84                 output = check_output_info[i];
85             }
86         }
87         return output;
88     }
89 }
```

## IV. Experiment Result

# Experiment Setting

- network architecture: 4-ary 3-cube
- flow control: 4 VCs per physical channel
- traffic patterns: *uniform random*, *torus tornado*, *torus transpose*, *shuffle*, *bit reverse* and *bit rotation*.
- compared topologies: Ring, Mesh\_XY
- sim-cycles: 20000

# Baselines

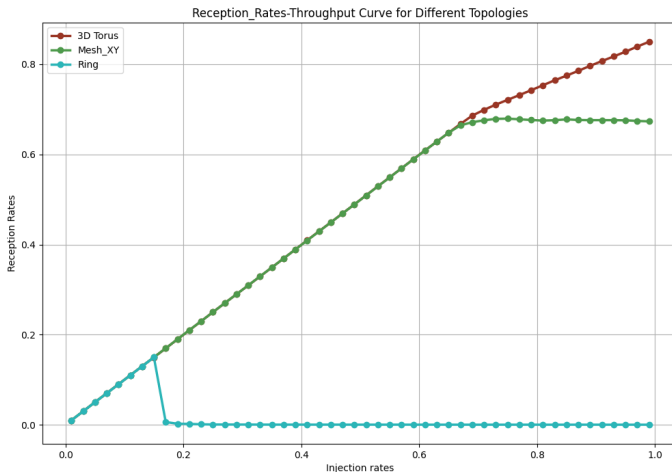
We use the following routing algorithms as baselines:

## baseline

- *default TABLE* is a minimal routing algorithm that is implemented in gem5. It does not include any flow control.
- *Dor* is a minimal routing algorithm that routes packets along the shortest path from dimension X to Z. We expand 2 VCs for each physical channel.
- *Minimal Adaptive* first decides the closest quadrant to the destination and then adaptively routes the packet within the quadrant. We expand 3 VCs for each physical channel.

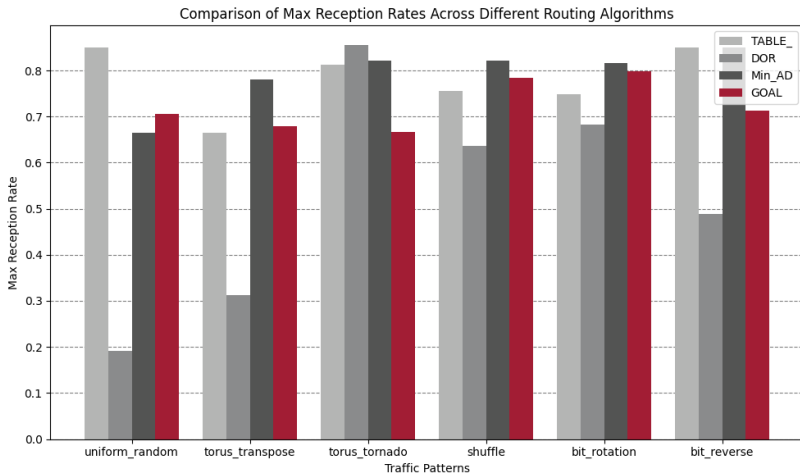
# Topology Comparison under Uniform Random Traffic

In this page we show comparisons on reception rates of three topologies.





# Max Reception Rate on all traffic patterns

In this page we show the saturation throughput of the four algorithms on a 4-ary 3-cube torus network on each traffic pattern.



# Average Latency on Injection Rate

In this page we show the latency changing with injection rate of the four algorithms under *bit rotation* and *torus tornado* respectively.







# Reception Rate on Injection Rate

In this page we show reception rate changing with injection rate under *uniform random* and *torus transpose* respectively.

It can be seen that GOAL is relatively robust under various traffic patterns and has a satisfying lower bound empirically.



# Conclusion

- The torus topology has a better performance than the ring and mesh topology because of its high path diversity, specifically, it has lower hops, higher and more stable throughput.
- The GOAL algorithm can balance the load among the quadrants selection and reduce the congestion in the network, which makes it perform more stable than other algorithms when the injection rates are increasing. But correspondingly, this algorithm sacrifices the low load performance.

# Division of Labor

- **Ziren Wang:** Implementation of the Torus topology, GOAL routing algorithm, FlowControl and the evaluation of the performance of the Torus topology and GOAL. Chapter 4- of report. PPT presentation.
- **Chuan Liu:** Implementation of Dimension Order Routing (Dor) baseline. the evaluation of the performance of the Torus topology and GOAL. Chapter 1-3 of report. PPT producing and PPT presentation.

# References

- Pablo E Berman, Luis Gravano, Gustavo D Pifarre, and Jorge LC Sanz. 1992. Adaptive deadlock-and livelock-free routing with all minimal paths in torus networks. In *Proceedings of the fourth annual ACM symposium on Parallel algorithms and architectures*. 3-12.
- Natalie Enright Jerger, Tushar Krishna, and Li-Shiuan Peh. 2022. *On-chip networks*. Springer Nature.
- Arjun Singh, William J Dally, Amit K Gupta, and Brian Towles. 2003. GOAL: a load balanced adaptive routing algorithm for torus networks. *ACM SIGARCH Computer Architecture News* 31, 2 (2003), 194-205.
- Leslie G. Valiant. 1982. A scheme for fast parallel communication. *SIAM journal on computing* 11, 2 (1982), 350-361.